

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
```

```
df=pd.read_csv('/content/Video_Games.csv')
df
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales	Critic_Score
0	Wii Sports	Wii	2006.0	Sports	Nintendo	41.36	28.96	3.77	8.45	82.53	
1	Super Mario Bros.	NES	1985.0	Platform	Nintendo	29.08	3.58	6.81	0.77	40.24	
2	Mario Kart Wii	Wii	2008.0	Racing	Nintendo	15.68	12.76	3.79	3.29	35.52	
3	Wii Sports Resort	Wii	2009.0	Sports	Nintendo	15.61	10.93	3.28	2.95	32.77	
4	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	Nintendo	11.27	8.89	10.22	1.00	31.37	
...	
16714	Samurai Warriors: Sanada Maru	PS3	2016.0	Action	Tecmo Koei	0.00	0.00	0.01	0.00	0.01	
16715	LMA Manager 2007	X360	2006.0	Sports	Codemasters	0.00	0.01	0.00	0.00	0.01	

```
df.shape
```

```
(16719, 16)
```

```
df.dtypes
```

```
Name          object
Platform       object
Year_of_Release  float64
Genre          object
Publisher       object
NA_Sales       float64
EU_Sales       float64
JP_Sales       float64
Other_Sales    float64
Global_Sales   float64
Critic_Score   float64
Critic_Count   float64
User_Score     object
User_Count     float64
Developer      object
Rating         object
dtype: object
```

```
df['Developer']=df['Developer'].str.lower()
df
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales
0	Wii Sports	Wii	2006.0	Sports	Nintendo	41.36	28.96
1	Super Mario Bros.	NES	1985.0	Platform	Nintendo	29.08	3.58
2	Mario Kart Wii	Wii	2008.0	Racing	Nintendo	15.68	12.76
3	Wii Sports Resort	Wii	2009.0	Sports	Nintendo	15.61	10.93
4	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	Nintendo	11.27	8.89
...
16714	Samurai Warriors: Sanada Maru	PS3	2016.0	Action	Tecmo Koei	0.00	0.00
16715	LMA Manager 2007	X360	2006.0	Sports	Codemasters	0.00	0.01

```
df['EU_Sales']=df['EU_Sales'].round(1)
df
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales
0	Wii Sports	Wii	2006.0	Sports	Nintendo	41.36	29.0
1	Super Mario Bros.	NES	1985.0	Platform	Nintendo	29.08	3.6
2	Mario Kart Wii	Wii	2008.0	Racing	Nintendo	15.68	12.8
3	Wii Sports Resort	Wii	2009.0	Sports	Nintendo	15.61	10.9
4	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	Nintendo	11.27	8.9
...
16714	Samurai Warriors: Sanada Maru	PS3	2016.0	Action	Tecmo Koei	0.00	0.0

The ceil function returns the smallest integer value which is greater than or equal to the specified number, whereas the floor function returns the largest integer value which is less than or equal to the specified number.

```
df['Other_Sales']=df['Other_Sales'].apply(np.ceil)
df
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales	EU_Sales
0	Wii Sports	Wii	2006.0	Sports	Nintendo	41.36	29.0
1	Super Mario Bros.	NES	1985.0	Platform	Nintendo	29.08	4.0
2	Mario Kart Wii	Wii	2008.0	Racing	Nintendo	15.68	13.0
3	Wii Sports Resort	Wii	2009.0	Sports	Nintendo	15.61	11.0
4	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	Nintendo	11.27	9.0
...
16714	Samurai Warriors: Sanada Maru	PS3	2016.0	Action	Tecmo Koei	0.00	0.0
16715	LMA Manager 2007	X360	2006.0	Sports	Codemasters	0.00	0.0

```
df['Other_Sales'].unique()
```

```
array([8.450e+00, 7.700e-01, 3.290e+00, 2.950e+00, 1.000e+00, 5.800e-01,
       2.880e+00, 2.840e+00, 2.240e+00, 4.700e-01, 2.740e+00, 1.900e+00,
       7.100e-01, 2.150e+00, 1.690e+00, 1.770e+00, 3.960e+00, 1.057e+01,
       5.500e-01, 2.040e+00, 1.360e+00, 4.200e-01, 4.600e-01, 1.410e+00,
       1.780e+00, 5.000e-01, 1.180e+00, 8.000e-01, 1.160e+00, 1.320e+00,
       5.900e-01, 2.380e+00, 1.130e+00, 7.800e-01, 2.420e+00, 1.120e+00,
       1.280e+00, 1.570e+00, 1.300e+00, 1.010e+00, 9.100e-01, 1.790e+00,
       1.970e+00, 8.600e-01, 1.210e+00, 2.300e-01, 7.600e-01, 7.400e-01,
       7.530e+00, 2.900e-01, 1.030e+00, 5.200e-01, 2.110e+00, 1.600e+00,
       1.610e+00, 3.500e-01, 9.700e-01, 1.060e+00, 6.300e-01, 1.500e-01,
       7.900e-01, 9.600e-01, 1.250e+00, 9.000e-01, 8.100e-01, 3.900e-01,
       6.800e-01, 8.500e-01, 1.800e-01, 8.000e-02, 6.700e-01, 7.000e-01,
       4.100e-01, 3.300e-01, 6.000e-01, 5.400e-01, 1.730e+00, 1.230e+00,
       1.600e-01, 1.110e+00, 3.100e-01, 4.800e-01, 6.200e-01, 1.900e-01,
       6.900e-01, 1.020e+00, 7.300e-01, 1.080e+00, 4.500e-01, 2.800e-01,
       5.100e-01, 2.200e-01, 1.090e+00, 9.900e-01, 3.000e-01, 6.400e-01,
       6.600e-01, 9.800e-01, 1.390e+00, 1.400e-01, 1.370e+00, 7.000e-02,
       2.100e-01, 6.100e-01, 1.700e-01, 1.200e-01, 0.000e+00, 7.200e-01,
       2.400e-01, 8.200e-01, 1.740e+00, 8.700e-01, 9.200e-01, 5.700e-01,
       1.100e-01, 4.000e-02, 5.600e-01, 2.000e-01, 3.400e-01, 9.000e-02,
       8.300e-01, 4.400e-01, 6.000e-02, 3.200e-01, 3.800e-01, 1.480e+00,
       3.700e-01, 1.000e-01, 2.500e-01, 3.600e-01, 1.300e-01, 4.300e-01,
       5.000e-02, 2.000e-02, 2.600e-01, 4.000e-01, 7.500e-01, 1.930e+00,
       8.400e-01, 5.300e-01, 8.900e-01, 1.670e+00, 2.700e-01, 2.930e+00,
       4.900e-01, 1.000e-02, 2.460e+00, 3.000e-02, 1.510e+00, 2.050e+00,
       1.680e+00, 1.820e+00, 1.330e+00, 9.400e-01, 9.300e-01])
```

```
df['Other_Sales'].isna()
```

```
0      False
1      False
2      False
3      False
4      False
...
16714   False
16715   False
16716   False
16717   False
16718   False
Name: Other_Sales, Length: 16719, dtype: bool
```

```
df['Other_Sales'].nunique()
```

```
155
```

```
missing_values=df.isnull().sum()
print(missing_values)
```

```
Name      2
Platform  0
Year_of_Release  269
Genre      2
Publisher  54
NA_Sales   0
```

```

EU_Sales      0
JP_Sales      0
Other_Sales    0
Global_Sales  0
Critic_Score   8582
Critic_Count   8582
User_Score     6704
User_Count     9129
Developer     6623
Rating         6769
dtype: int64

```

```

data=pd.read_csv('/content/homelessness (1).csv')
data
print(data.head(10))

```

```

   Unnamed: 0  region  state  individuals \
0           0  East South Central  Alabama    2570.0
1           1         Pacific  Alaska    1434.0
2           2         Mountain  Arizona    7259.0
3           3  West South Central  Arkansas    2280.0
4           4         Pacific  California  109008.0
5           5         Mountain  Colorado    7607.0
6           6     New England  Connecticut    2280.0
7           7  South Atlantic  Delaware     708.0
8           8  South Atlantic  District of Columbia    3770.0
9           9  South Atlantic  Florida   21443.0

   family_members  state_pop
0           864.0    4887681
1           582.0    735139
2          2606.0   7158024
3           432.0   3009733
4         20964.0  39461588
5          3250.0   5691287
6          1696.0   3571520
7           374.0   965479
8          3134.0   701547
9          9587.0  21244317

```

```

data.dropna(axis=1,inplace=True)
data

```

	Unnamed: 0	region	state	individuals	family_members	state_pop	family
0	0	East South Central	Alabama	2570.0	0.02	4887681	0.016593
1	1	Pacific	Alaska	1434.0	0.01	735139	0.011177
2	2	Mountain	Arizona	7259.0	0.05	7158024	0.050048
3	3	West South Central	Arkansas	2280.0	0.01	3009733	0.008297
4	4	Pacific	California	109008.0	0.40	39461588	0.402612
5	5	Mountain	Colorado	7607.0	0.06	5691287	0.062416
6	6	New England	Connecticut	2280.0	0.03	3571520	0.032572
7	7	South Atlantic	Delaware	708.0	0.01	965479	0.007183
8	8	South Atlantic	District of Columbia	3770.0	0.06	701547	0.060188
9	9	South Atlantic	Florida	21443.0	0.18	21244317	0.184118
10	10	South Atlantic	Georgia	6943.0	0.05	10511131	0.049088
11	11	Pacific	Hawaii	4131.0	0.05	1420593	0.046073
12	12	Mountain	Idaho	1297.0	0.01	1750536	0.013732
13	13	East North Central	Illinois	6752.0	0.07	12723071	0.074726

▼ single feature scaling

It converts every value of a column between 0 and 1. Every number of a column will be divided by the max no of a column.

```
13      13      130000      10000      1711.0      0.02      3140010      0.013933
```

```
data['family']=data['family_members']/data['family_members'].max()
data.head()
```

	Unnamed: 0	region	state	individuals	family_members	state_pop	family
0	0	East South Central	Alabama	2570.0	0.016593	4887681	0.016593
1	1	Pacific	Alaska	1434.0	0.011177	735139	0.011177
2	2	Mountain	Arizona	7259.0	0.050048	7158024	0.050048
3	3	West South Central	Arkansas	2280.0	0.008297	3009733	0.008297

```
data['family_members']=data['family'].round(2)
data.head()
```

	Unnamed: 0	region	state	individuals	family_members	state_pop	family
0	0	East South Central	Alabama	2570.0	0.02	4887681	0.016593
1	1	Pacific	Alaska	1434.0	0.01	735139	0.011177
2	2	Mountain	Arizona	7259.0	0.05	7158024	0.050048
3	3	West South Central	Arkansas	2280.0	0.01	3009733	0.008297

▼ Min-Max

It same as single feature scaling, here also value of a column between 0 and 1. Every no.of coulumn will be subtractedby the min no.of that column and divided by the range of that column.

```
data['state_pop']=(data['state_pop']-data['state_pop'].min())/(data['state_pop'].max()-data['state_pop'].min())
data.head()
```

	Unnamed: 0	region	state	individuals	family_members	state_pop	family
0	0	East South Central	0.110845	2570.0	0.02	0.110845	0.016593
1	1	Pacific	0.004051	1434.0	0.01	0.004051	0.011177
2	2	Mountain	0.169232	7259.0	0.05	0.169232	0.050048
3	3	West South Central	0.062548	2280.0	0.01	0.062548	0.008297

▼ Zscore

It converts every value of a column into a number around 0. Typical values obtained by a z-score transformation range from -3 and 3. The new value is calculated as the difference between the current value and the average value, divided by the standard deviation. The average value of a column can be obtained through the `mean()` function, while the standard deviation through the `std()` function.

```
data['individuals']=(data['individuals']-data['individuals'].mean())/data['individuals'].std()
data.head()
```

▼ Data Normalization technique on categorical data

We used the label encoder library file and some variables to store the columns and transform to fit them.

```
col_to_encode = 'region'
data_to_encode = data[col_to_encode]
lae = LabelEncoder()
lae.fit(data_to_encode)

encoded_data = lae.transform(data_to_encode)
data[col_to_encode] = encoded_data
data.head()
```

	Unnamed: 0	region	state	individuals	family_members	state_pop
0	0	1	Alabama	2570.0	864.0	4887681
1	1	5	Alaska	1434.0	582.0	735139
2	2	3	Arizona	7259.0	2606.0	7158024
3	3	8	Arkansas	2280.0	432.0	3009733
4	4	5	California	109008.0	20964.0	39461588

A lambda function is a small anonymous function. A lambda function can take any number of arguments, but can only have one expression.

```
def myfunc(n):
    return lambda a : a * n

mytripler = myfunc(2)

print(mytripler(11))
```

22

apply() must be applied to every row (through the parameter axis = 1) and then through the lambda operator we can select the specific row and apply it the function set_pattern().

```
import re
def set_pattern(x):
    pattern = r'[(A-Z)]\w+,[([A-Z])]\w+'
    res = re.match(pattern, x)
    if res:
        x = x.replace(',', ' ', ' ')
    return x
```

```
data['state'] = data.apply(lambda x: set_pattern(x['state']), axis=1)
print(data)
```


Unnamed: 0	region	state	individuals	family_members	\
0	0	1	Alabama	2570.0	864.0
1	1	5	Alaska	1434.0	582.0
2	2	3	Arizona	7259.0	2606.0
3	3	8	Arkansas	2280.0	432.0
4	4	5	California	109008.0	20964.0
5	5	3	Colorado	7607.0	3250.0
6	6	4	Connecticut	2280.0	1696.0
7	7	6	Delaware	708.0	374.0
8	8	6	District of Columbia	3770.0	3134.0
9	9	6	Florida	21443.0	9587.0
10	10	6	Georgia	6943.0	2556.0
11	11	5	Hawaii	4131.0	2399.0
12	12	3	Idaho	1297.0	715.0
13	13	0	Illinois	6752.0	3891.0
14	14	0	Indiana	3776.0	1482.0
15	15	7	Iowa	1711.0	1038.0
16	16	7	Kansas	1443.0	773.0
17	17	1	Kentucky	2735.0	953.0
18	18	8	Louisiana	2540.0	519.0
19	19	4	Maine	1450.0	1066.0
20	20	6	Maryland	4914.0	2230.0
21	21	4	Massachusetts	6811.0	13257.0
22	22	0	Michigan	5209.0	3142.0
23	23	7	Minnesota	3993.0	3250.0
24	24	1	Mississippi	1024.0	328.0
25	25	7	Missouri	3776.0	2107.0
26	26	3	Montana	983.0	422.0
27	27	7	Nebraska	1745.0	676.0
28	28	3	Nevada	7058.0	486.0
29	29	4	New Hampshire	835.0	615.0
30	30	2	New Jersey	6048.0	3350.0
31	31	3	New Mexico	1949.0	602.0
32	32	2	New York	39827.0	52070.0
33	33	6	North Carolina	6451.0	2817.0
34	34	7	North Dakota	467.0	75.0
35	35	0	Ohio	6929.0	3320.0
36	36	8	Oklahoma	2823.0	1048.0
37	37	5	Oregon	11139.0	3337.0
38	38	2	Pennsylvania	8163.0	5349.0
39	39	4	Rhode Island	747.0	354.0
40	40	6	South Carolina	3082.0	851.0
41	41	7	South Dakota	836.0	323.0
42	42	1	Tennessee	6139.0	1744.0
43	43	8	Texas	19199.0	6111.0
44	44	3	Utah	1904.0	972.0
45	45	4	Vermont	780.0	511.0
46	46	6	Virginia	3928.0	2047.0
47	47	5	Washington	16424.0	5880.0
48	48	6	West Virginia	1021.0	222.0
49	49	0	Wisconsin	2740.0	2167.0
50	50	3	Wyoming	434.0	205.0

state_pop	
0	4887681
1	735139
2	7158024
3	3009733

▼ Data Analysis

1. Univariate Analysis: Count Plot, Histogram, Skewness, Kurtosis, Joint Plot, Box Plot, Pie Chart, Line Chart, Distribution Plot
2. Bivariate Analysis: Scatter Plot, Joint Plot, Line Plot, Bar Plot, Box Plot
3. Multivariate Analysis: Pair Plot, Heatmap, Bar Plot, Line Plot, Scatter Plot

IV. Perform the following data exploration analysis on the downloaded dataset:

1. Count Plot
2. Box Plot
3. Joint Plot
4. Correlation
5. Skewness
6. Distribution Plot
7. Pair Plot
8. Heatmap


```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style()
import numpy as pd
```

```
den=pd.read_csv("/content/tweets.csv")
```

```
den.head()
```

	Tweet Id	Tweet URL	Tweet Posted Time (UTC)	Tweet Content
0	"1167429261210218497"	https://twitter.com/animalhealthEU/status/1167...	30 Aug 2019 13:30:00	Pets change lives & become a p
1	"1167375334670557185"	https://twitter.com/PennyBrohnUK/status/116737...	30 Aug 2019 09:55:43	Another spot o #morethanmed bus ir
2	"1167237977615097861"	https://twitter.com/lordbyronaf/status/1167237...	30 Aug 2019 00:49:54	What a great t @HealthSourc @Local
3	"1167236897078480898"	https://twitter.com/CountessDavis/status/11672...	30 Aug 2019 00:45:37	What a great t @HealthSourc @Local
4	"1167228378191204353"	https://twitter.com/Local12/status/11672283781...	30 Aug 2019 00:11:46	What a great t @HealthSourc @Local


5 rows × 21 columns



▼ count plot

The countplot is used to represent the occurrence(counts) of the observation present in the variable. It uses the concept of a bar chart for the visual depiction.


```
sns.countplot(x='Tweet Type',data=den)  
plt.show()
```




▼ Box plot

Box Plot is also known as Whisker plot is created to display the summary of the set of data values having properties like minimum, first quartile, median, third quartile and maximum.

```
sns.boxplot(x='Likes Received',data=den)  
plt.show()
```




```
sns.boxplot(x='User Followers',data=den)  
plt.show()
```




▼ Joint plot

The joint plot is a way of understanding the relationship between two variables and the distribution of individuals of each variable. kind : { "scatter" | "reg" | "resid" | "kde" | "hex" } Kind of plot to draw.


```
sns.jointplot(x='User Followers',y='Likes Received',data=den, kind='reg')
plt.show()
```




```
sns.jointplot(x='User Followers',y='Likes Received',data=den, kind='scatter')
plt.show()
```



```
sns.jointplot(x='User Followers',y='Likes Received',data=den, kind='kde')  
plt.show()
```



```
sns.jointplot(x='User Followers',y='Likes Received',data=den, kind='hex')  
plt.show()
```



▼ correlation

The `pearsonr()` SciPy function can be used to calculate the Pearson's correlation coefficient between two data samples with the same length. Correlation is a statistical measure that expresses the extent to which two variables are linearly related. A null hypothesis is a type of statistical hypothesis that proposes that no statistical significance exists in a set of given observations.

```
from scipy.stats import pearsonr

def get_correlation(column1, column2, df):
    pearson_corr, p_value = pearsonr(df[column1], df[column2])
    print("Correlation between {} and {} is {}".format(column1, column2, pearson_corr))
    print("P-value of this correlation is {}".format(p_value))
```

```
get_correlation('User Followers', 'Likes Received', den)
```

```
Correlation between User Followers and Likes Received is 0.005301145045724385
P-value of this correlation is 0.9173170162772105
```

▼ skewness

Skewness is a measure of the asymmetry of a distribution. A distribution is asymmetrical when its left and right side are not mirror images.

```
from scipy.stats import skew
den.skew(axis = 0, skipna = True)
plt.show()
```

```
<ipython-input-17-bb391122d770>:2: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select or
den.skew(axis = 0, skipna = True)
```

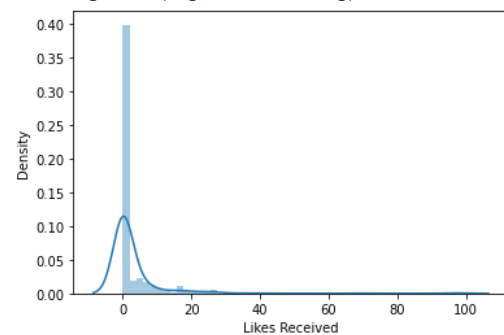
```
den['Likes Received'].skew(axis = 0, skipna = True)
plt.show()
```

▼ Distribution plot

`distplot()` function is used to plot the distplot. The distplot represents the univariate distribution of data. The `seaborn.distplot()` function accepts the data variable as an argument and returns the plot with the density distribution.

```
sns.distplot(den['Likes Received'])  
plt.show()
```


```
/usr/local/lib/python3.9/dist-packages/seaborn/distributions.py:2619: FutureWarning: `dist`  
warnings.warn(msg, FutureWarning)
```



▼ Pair plot

The Seaborn Pairplot allows us to plot pairwise relationships between variables within a dataset. This creates a nice visualisation and helps us understand the data by summarising a large amount of data in a single figure.


```
sns.pairplot(den)  
plt.show()
```



Heat map

A heatmap is a two-dimensional graphical representation of data where the individual values that are contained in a matrix are represented as colours. The Seaborn package allows the creation of annotated heatmaps which can be tweaked using Matplotlib tools as per the creator's requirement.

```
sns.heatmap(den.corr(),annot=True,fmt='.1f',linewidths=2)
plt.show()
```



5. Frame a problem statement based on the attributes of the dataset and predict the accuracy of the dependent variable using the values of the independent variables by using the following machine learning models:'

-> Logistic Regression Logistic regression aims to solve classification problems. It does this by predicting categorical outcomes, unlike linear regression that predicts a continuous outcome.

-> Linear SVC Linear Support Vector Machine (Linear SVC) is an algorithm that attempts to find a hyperplane to maximize the distance between classified samples.

-> Decision Tree Classifier Decision Trees (DTs) are a non-parametric supervised learning method used for classification and regression. The goal is to create a model that predicts the value of a target variable by learning simple decision rules inferred from the data features. A tree can be seen as a piecewise constant approximation.

-> Naive Bayes - Gaussian NB Gaussian Naïve Bayes classifier assumes that the data from each label is drawn from a simple Gaussian distribution. The Scikit-learn provides `sklearn.naive_bayes.GaussianNB` to implement the Gaussian Naïve Bayes algorithm for classification.

-> Ensemble Model - RandomForestClassifier, GradientBoostingClassifier, AdaBoostClassifier

```
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
df=pd.read_csv('/content/archive (3).zip')
df
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 9 columns

```
df=df[['BMI', 'Glucose', 'BloodPressure', 'Insulin', 'Age']]
df
```

	BMI	Glucose	BloodPressure	Insulin	Age
0	33.6	148	72	0	50
1	26.6	85	66	0	31
2	23.3	183	64	0	32
3	28.1	89	66	94	21
4	43.1	137	40	168	33

```
df['Glucose'].fillna(df['Glucose'].median())
```

```
0    148
1     85
2    183
3     89
4    137
```

```
...
763   101
764   122
765   121
766   126
767    93
```

```
Name: Glucose, Length: 768, dtype: int64
```

```
x=df.drop('Glucose',axis=1)
y=df['Glucose']
```

```
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

```
# load the diabetes dataset from a CSV file
diabetes_df = pd.read_csv('/content/diabetes.csv')
diabetes_df
```

Pregnancies Glucose BloodPressure SkinThickness Insulin BMI DiabetesPedigreeFu

```
# split the dataset into features (X) and target (y)
X = diabetes_df.drop('Outcome', axis=1)
y = diabetes_df['Outcome']
```

```
# split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
# create a logistic regression model
logreg = LogisticRegression()
```

```
764      2      122      70      27      0  36.8
```

```
# fit the model to the training data
logreg.fit(X_train, y_train)
```

```
# predict the target values for the testing data
y_pred = logreg.predict(X_test)
```

```
# evaluate the performance of the model
score = logreg.score(X_test, y_test)
```

```
print("Accuracy:", score)
```

```
Accuracy: 0.7467532467532467
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
from sklearn.svm import LinearSVC
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

```
# split the dataset into features (X) and target (y)
X = diabetes_df.drop('Outcome', axis=1)
y = diabetes_df['Outcome']
```

```
# split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
# create a linear support vector classification model
svc = LinearSVC()
```

```
# fit the model to the training data
svc.fit(X_train, y_train)
```

```
# predict the target values for the testing data
y_pred = svc.predict(X_test)
```

```
# evaluate the performance of the model
```

```
score = accuracy_score(y_test, y_pred)
```

```
print("Accuracy:", score)
```

```
Accuracy: 0.6363636363636364  
/usr/local/lib/python3.9/dist-packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.  
  warnings.warn(
```

```
from sklearn.tree import DecisionTreeClassifier  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
# create a decision tree classification model  
dtc = DecisionTreeClassifier()
```

```
# fit the model to the training data  
dtc.fit(X_train, y_train)
```

```
# predict the target values for the testing data  
y_pred = dtc.predict(X_test)
```

```
# evaluate the performance of the model  
score = accuracy_score(y_test, y_pred)
```

```
print("Accuracy:", score)
```

```
Accuracy: 0.7142857142857143
```

```
from sklearn.naive_bayes import GaussianNB  
# create a Gaussian Naive Bayes classifier  
gnb = GaussianNB()
```

```
# fit the model to the training data  
gnb.fit(X_train, y_train)
```

```
# predict the target values for the testing data  
y_pred = gnb.predict(X_test)
```

```
# evaluate the performance of the model  
score = accuracy_score(y_test, y_pred)
```

```
print("Accuracy:", score)
```

```
Accuracy: 0.7662337662337663
```

```
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, AdaBoostClassifier  
# create a Random Forest classifier  
rfc = RandomForestClassifier()
```

```
# fit the model to the training data  
rfc.fit(X_train, y_train)
```

```
# predict the target values for the testing data
```

```
y_pred = rfc.predict(X_test)

# evaluate the performance of the model
score = accuracy_score(y_test, y_pred)

print("Random Forest Classifier Accuracy:", score)

# create a Gradient Boosting classifier
gbc = GradientBoostingClassifier()

# fit the model to the training data
gbc.fit(X_train, y_train)

# predict the target values for the testing data
y_pred = gbc.predict(X_test)

# evaluate the performance of the model
score = accuracy_score(y_test, y_pred)

print("Gradient Boosting Classifier Accuracy:", score)

# create an AdaBoost classifier
abc = AdaBoostClassifier()

# fit the model to the training data
abc.fit(X_train, y_train)

# predict the target values for the testing data
y_pred = abc.predict(X_test)

# evaluate the performance of the model
score = accuracy_score(y_test, y_pred)

print("AdaBoost Classifier Accuracy:", score)

Random Forest Classifier Accuracy: 0.7727272727272727
Gradient Boosting Classifier Accuracy: 0.7662337662337663
AdaBoost Classifier Accuracy: 0.7727272727272727
```